

Introduction

Locksmith is a tool that performs static analysis in order to detect data races in C programs. In general, data races happen when two or more threads attempt to read on a specific memory address while at least one of them attempts to write on it. Thus, Locksmith tries to predict which memory addresses are prone to have data races by reading the program in compile time.

Locksmith

Locksmith's algorithm is divided in a number of phases. These phases are as follows:

Typing phase

Locksmith reads the entire program and stores its control flow. All information about the control flow is stored in phi (ϕ) structs. Information about the global variables is stored in rho (ρ) structs.

Shared phase

In this phase, Locksmith finds all variables that could be used simultaneously among the threads. These variables/memory addresses are the ones who need to be checked for data races. Other variables, such as thread local variables, are not prone to data races and thus are excluded from this analysis.

Linearity phase

Checks if the locks used are linear or non-linear. If a lock is non-linear, it means that it could represent multiple locks and Locksmith cannot know which lock is acquired in this case.

Lock State phase

Locksmith finds the state of each lock after each statement of the program i.e. if a lock is acquired at any point or not.

Guarded-By phase

In this phase, Locksmith checks each shared variable/memory address which locks are held each time they are accessed. More specifically, it initializes guards in each dereference of shared variables which stores a correlation of a variable with a corresponding lock, if any. Then, for each guard, it propagates through the program flow backwards until it reaches the main() function, possibly creating more correlations and guards on its way. It ends once every correlation has propagated to the main() function. After this phase, Locksmith knows where each variable acquires or releases a lock at any point, so it knows all the potential data races.

Escapes phase

Here, Locksmith completes the work done in Linearity Phase.

Races phase

Print whether a variable could have a potential data race as a warning, along with that variable's references.

Locksmith output

Now, consider the following program:

test.c

```
#include <pthread.h>

pthread_mutex_t lock1, lock2, lock[2];
float shared_var;
int partly_unprotected;
float non_linear;
int many_refs;
int completely_unprotected;

void *my_func(void *arg) {
    int local_var = 4, local_var2;
    pthread_mutex_lock(&lock1);
    shared_var++;
    many_refs += 5;
    pthread_mutex_unlock(&lock1);
    pthread_mutex_lock(&lock2);
    partly_unprotected = local_var;
    local_var += many_refs;
    pthread_mutex_unlock(&lock2);
    many_refs = 1;
    local_var = partly_unprotected;
    completely_unprotected--;
    pthread_mutex_lock(&lock[1]);
    non_linear = 4.2;
    pthread_mutex_unlock(&lock[1]);
    local_var2 = many_refs;
    return NULL;
}

int main() {
    pthread_t t1, t2, t[2];

    int i;
    for(i = 0; i < 2; i++)
```

```

{
    pthread_mutex_init(&lock[i], NULL);
    pthread_create(&t[i], NULL, my_func, NULL);
}

pthread_mutex_init(&lock1, NULL);
pthread_mutex_init(&lock2, NULL);

pthread_create(&t1, NULL, my_func, NULL);
pthread_create(&t2, NULL, my_func, NULL);

return 1;
}

```

This is Locksmith's output for this program:

Warning: Possible data race: &partly_unprotected:test.c:5 is not protected!
references:

dereference of &partly_unprotected:test.c:5 at test.c:17

&partly_unprotected:test.c:5

locks acquired:

concrete lock2:test.c:41

lock2:test.c:3

in: main at test.c:30 -> test.c:37

in: main at test.c:30 -> test.c:43

in: main at test.c:30 -> test.c:44

dereference of &partly_unprotected:test.c:5 at test.c:21

&partly_unprotected:test.c:5

locks acquired:

<empty>

in: main at test.c:30 -> test.c:37

in: main at test.c:30 -> test.c:43

in: main at test.c:30 -> test.c:44

Warning: Possible data race: &non_linear:test.c:6 is protected by non-linear or concrete lock(s):

*lock:test.c:3

non linear concrete lock[i]:test.c:36

references:

dereference of &non_linear:test.c:6 at test.c:24

&non_linear:test.c:6

locks acquired:
 *lock:test.c:3
 non linear concrete lock[i]:test.c:36
in: main at test.c:30 -> test.c:37
in: main at test.c:30 -> test.c:43
in: main at test.c:30 -> test.c:44

Warning: Possible data race: &many_refs:test.c:7 is not protected!

references:
dereference of &many_refs:test.c:7 at test.c:14

 &many_refs:test.c:7

locks acquired:
 concrete lock1:test.c:40
 lock1:test.c:3
in: main at test.c:30 -> test.c:37
in: main at test.c:30 -> test.c:43
in: main at test.c:30 -> test.c:44

dereference of &many_refs:test.c:7 at test.c:18

 &many_refs:test.c:7

locks acquired:
 concrete lock2:test.c:41
 lock2:test.c:3
in: main at test.c:30 -> test.c:37
in: main at test.c:30 -> test.c:43
in: main at test.c:30 -> test.c:44

dereference of &many_refs:test.c:7 at test.c:20

 &many_refs:test.c:7

locks acquired:
 <empty>
in: main at test.c:30 -> test.c:37
in: main at test.c:30 -> test.c:43
in: main at test.c:30 -> test.c:44

dereference of &many_refs:test.c:7 at test.c:26

 &many_refs:test.c:7

locks acquired:
 <empty>
in: main at test.c:30 -> test.c:37
in: main at test.c:30 -> test.c:43
in: main at test.c:30 -> test.c:44

```
Warning: Possible data race: &completely_unprotected:test.c:8 is not protected!
references:
  dereference of &completely_unprotected:test.c:8 at test.c:22
    &completely_unprotected:test.c:8
locks acquired:
  <empty>
in: main at test.c:30 -> test.c:37
in: main at test.c:30 -> test.c:43
in: main at test.c:30 -> test.c:44
```

As we can see, of all the variables used, `many_refs`, `completely_unprotected` and `partly_unprotected` are not protected and thus prone to data races. In the case of `completely_unprotected`, it is used twice (decrement both reads and writes a value) and no lock was ever acquired (and released) at this point. The `partly_unprotected` variable has two references: in the first one it acquired one lock before being accessed but in the second one it didn't do so. In the case of `many_refs`, it is used in 4 statements but is only protected in two of them. `Non-linear` is only accessed once but it uses a non-linear lock. As for the other variables, `shared_var` is accessed once and has got a lock at that point, so it is protected. `local_var` and `local_var2` are thread local variables, so they are not shared with the other thread and there is no need to check if they are protected. Notice that the warnings are shown in the order their respective variables were first declared. But what happens if we have multiple warnings? How do we know which warnings are more important than the others?

Evaluation Criteria

Warnings are sorted in descending order of importance. In order to sort the warnings, Locksmith uses certain evaluation criteria. These criteria are as follows:

The number of times a memory address is accessed for writing

As was already mentioned, in order for data races to exist, the threads have to write in a shared memory address. If there are multiple writing operations for a specific variable, it probably means that it needs more to be done to fix the data race than a race with fewer writing operations.

The number of total references of a memory address

Each access of a memory address constitutes a reference to it. Similarly with the first criterion, if a variable is accessed a lot in a program, it probably means that it is more impactful for the program than a variable that is accessed fewer times. In this case, a warning regarding the former variable will be considered more important than the latter.

Locks acquired by a variable

Locks prevent simultaneous access to a memory address by multiple threads. This usually prevents data races, however, locks should be acquired every time before a shared variable is accessed and released each time after it is done. It only needs to be accessed once without a lock for a data race to occur. While a data race would be equally disastrous for a program if a variable used no locks and another variable used locks in a few read/write operations, the latter is easier to fix and thus will be considered less important than the former.

A memory address is protected by non-linear locks

Sometimes a memory address may be protected by a lock. However, in some cases, different locks may be created in the same line in the program (e.g. inside a loop) and using these locks may cause a warning because the locks would be non-linear. Because Locksmith is unable to tell apart two function calls in the same line, it cannot possibly know whether it can cause an actual data race or not. As a result, those warnings are not deemed important.

Implementation

Each concrete location has a value called `rho_priority_value`. The higher that number is, the higher the importance of a warning regarding it. These locations are stored in the `all_concrete_rho` reference. Throughout Locksmith's analysis, it updates these values whenever necessary with a positive modifier depending on the evaluation criteria that increase the importance of the warnings. For example, a variable that is accessed once for writing (+20000) receives a larger modifier than a variable that is accessed once for reading (+10000). Similarly, it gives a negative modifier depending on the evaluation criteria that decrease the importance of the warnings. For example, a variable that acquired one lock (-10000) receives a smaller negative modifier than a variable with a non-linear lock (-1000000000). Once the analysis is done, it sorts `all_concrete_rho` in descending order and prints those who may cause warnings. It is important to note that `rho_priority_value` will be updated for all concrete locations during the analysis, regardless if it leads to a race or not.

Now, consider the previous program again. If we run Locksmith again by ranking the warnings (which does by default) we get this:

```
Warning: Possible data race: &many_refs:test.c:7 is not protected!  
references:  
  dereference of &many_refs:test.c:7 at test.c:14  
    &many_refs:test.c:7  
locks acquired:  
  concrete lock1:test.c:40  
  lock1:test.c:3
```

in: main at test.c:30 -> test.c:37
in: main at test.c:30 -> test.c:43
in: main at test.c:30 -> test.c:44

dereference of &many_refs:test.c:7 at test.c:18
 &many_refs:test.c:7

locks acquired:
 concrete lock2:test.c:41
 lock2:test.c:3

in: main at test.c:30 -> test.c:37
in: main at test.c:30 -> test.c:43
in: main at test.c:30 -> test.c:44

dereference of &many_refs:test.c:7 at test.c:20
 &many_refs:test.c:7

locks acquired:
 <empty>

in: main at test.c:30 -> test.c:37
in: main at test.c:30 -> test.c:43
in: main at test.c:30 -> test.c:44

dereference of &many_refs:test.c:7 at test.c:26
 &many_refs:test.c:7

locks acquired:
 <empty>

in: main at test.c:30 -> test.c:37
in: main at test.c:30 -> test.c:43
in: main at test.c:30 -> test.c:44

Warning: Possible data race: &completely_unprotected:test.c:8 is not protected!
references:

dereference of &completely_unprotected:test.c:8 at test.c:22
 &completely_unprotected:test.c:8

locks acquired:
 <empty>

in: main at test.c:30 -> test.c:37
in: main at test.c:30 -> test.c:43
in: main at test.c:30 -> test.c:44

Warning: Possible data race: &partly_unprotected:test.c:5 is not protected!
references:

dereference of &partly_unprotected:test.c:5 at test.c:17

```
&partly_unprotected:test.c:5
locks acquired:
  concrete lock2:test.c:41
  lock2:test.c:3
in: main at test.c:30 -> test.c:37
in: main at test.c:30 -> test.c:43
in: main at test.c:30 -> test.c:44

dereference of &partly_unprotected:test.c:5 at test.c:21
  &partly_unprotected:test.c:5
locks acquired:
  <empty>
in: main at test.c:30 -> test.c:37
in: main at test.c:30 -> test.c:43
in: main at test.c:30 -> test.c:44
```

Warning: Possible data race: &non_linear:test.c:6 is protected by non-linear or concrete lock(s):

```
*lock:test.c:3
non linear concrete lock[i]:test.c:36
references:
dereference of &non_linear:test.c:6 at test.c:24
  &non_linear:test.c:6
locks acquired:
  *lock:test.c:3
  non linear concrete lock[i]:test.c:36
in: main at test.c:30 -> test.c:37
in: main at test.c:30 -> test.c:43
in: main at test.c:30 -> test.c:44
```

This time, many_refs appears first. In total, it has 2 writing operations, 3 reading operations and acquired locks twice for a value of $2*20000 + 3*10000 - 2*10000 = 50000$. Next up is completely_unprotected with 1 writing operation and 1 reading operation for a value of $20000 + 10000 = 30000$. After that we have partly_unprotected with 1 writing operation, 1 reading operation and 1 lock acquired for a value of $20000 + 10000 - 10000 = 20000$. Finally, we have non_linear with 1 writing operation, 1 lock acquired and 1 non-linear lock for a value of $20000 - 10000 - 10000000000 = -9999900000$. As we can see, the warnings are indeed mentioned in descending order of their rho_priority_values.

Conclusion

So what do these criteria have in common? They grade the warnings depending on how much work needs to be done in order to be fixed. They assume that a programmer will run Locksmith for a program to check for data races, so Locksmith will need to first show the warnings that affect the program the most, with the least impactful warnings being shown last.